

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
on

Machine Learning (23CS6PCMAL)

Submitted by

AKANKSHA SINGA (1BM22CS027)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Sep-2024 to Jan-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Akanksha Singa(1BM22CS027)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

M Lakshmi Neelima Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
----------------------------------------------------------------------	------------------------------------------------------------------

Index

Sl. No.	Date	Experiment Title	Page No.
1	21-2-2025	Write a python program to import and export data using Pandas library functions	2-7
2	3-3-2025	Demonstrate various data pre-processing techniques for a given dataset	8-11
3	10-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	12-16
4	17-3-2025	Build Logistic Regression Model for a given dataset	17-21
5	24-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	22-24
6	7-4-2025	Build KNN Classification model for a given dataset.	25-28
7	21-4-2025	Build Support vector machine model for a given dataset	29-31
8	5-5-2025	Implement Random forest ensemble method on a given dataset.	32-33
9	5-5-2025	Implement Boosting ensemble method on a given dataset.	34-35
10	12-5-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	36-38
11	12-5-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	39-42

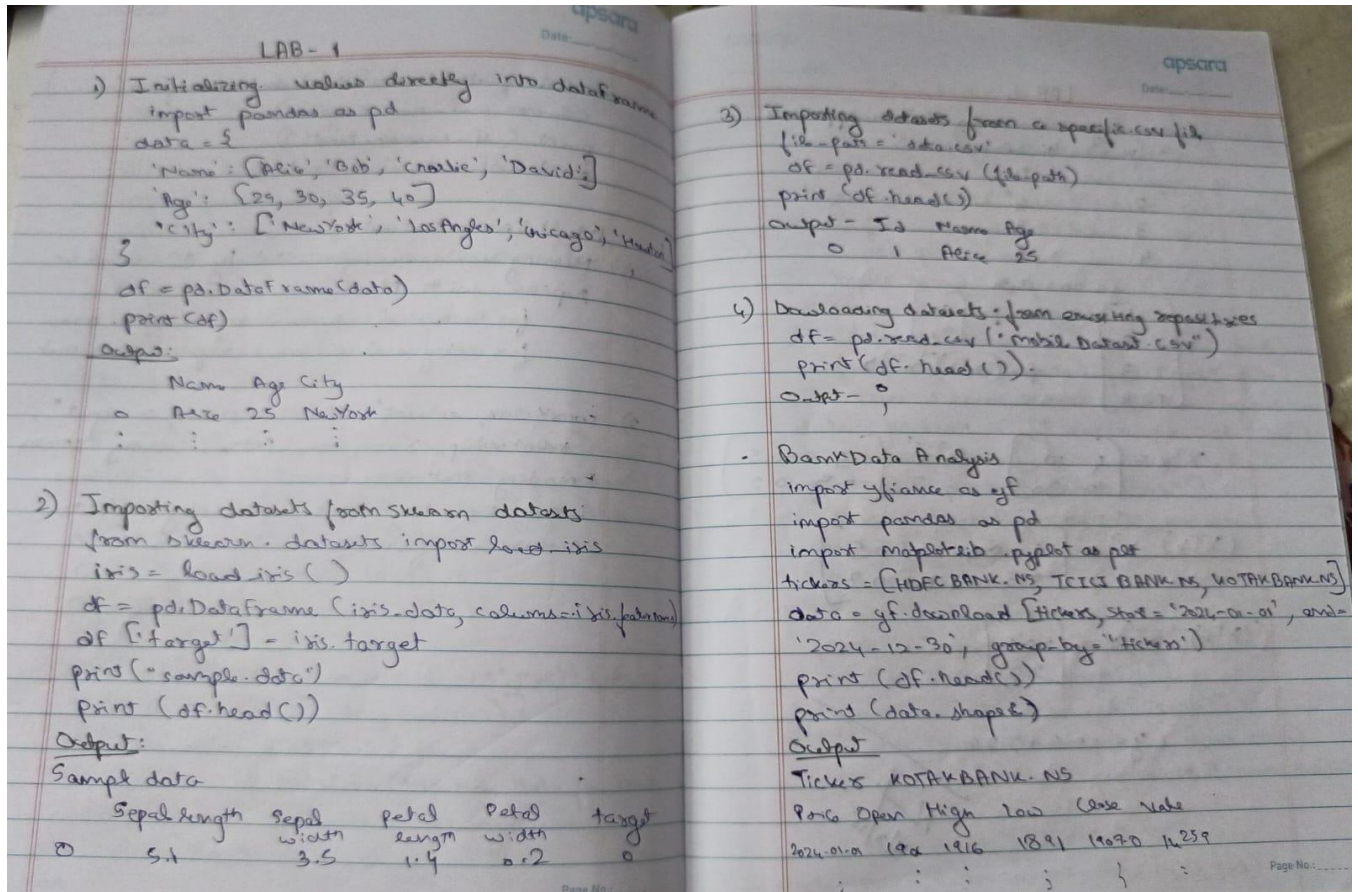
Github Link:

https://github.com/Akanksha-singa/6A_ML_LAB_B2

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot:



Code:

```
import pandas as pd
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [25, 30, 35, 40],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}
df = pd.DataFrame(data)
print("Sample data:")
print(df.head())
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("Sample data:")
```

```

print(df.head())
file_path = 'data.csv'
df = pd.read_csv(file_path)
print("Sample data:")
print(df.head())
print("\n")
file_path = 'mobiles-dataset-2025.csv'
df = pd.read_csv(file_path, encoding='latin-1') # or 'cp1252' or other suitable
encoding print("Sample data:")
print(df.head())
import pandas as pd

data = {
'USN': ['IS001','IS002','IS003','IS004','IS005'],
'Name': ['Alice', 'Bob', 'Charlie', 'David','Eve'],
'Marks': [25, 30, 35, 40,45]
}

df = pd.DataFrame(data)
print("Sample data:")
print(df.head())
from sklearn.datasets import load_diabetes
iris = load_diabetes()
df = pd.DataFrame(iris.data, columns=iris.feature_names)

print("Sample data:")
print(df.head())
file_path = 'sample_sales_data.csv'
df = pd.read_csv(file_path)
print("Sample data:")
print(df.head())
print("\n")

df = pd.read_csv("/content/dataset-of-diabetes .csv",encoding='latin-1')
print("Sample data:")
print(df.head())
print("\n")

df =pd.read_csv('sample_sales_data.csv')
print("Sample data:")
print(df.head())

df.to_csv('output.csv',index=False)
print("Data saved to output.csv")
sales_df =pd.read_csv('sample_sales_data.csv')
print("Sample data:")
print(sales_df.head())
sales_by_region =sales_df.groupby('Region')['Sales'].sum()
print("\nTotal sales by region:")
print(sales_by_region)
best_selling_products

```

```

=sales_df.groupby('Product')['Quantity'].sum().sort_values(ascending=False)
print("\nBest-selling products by quantity:")
print(best_selling_products)
sales_by_region.to_csv('sales_by_region.csv')
best_selling_products.to_csv('best_selling_products.csv')
print("Data saved to sales_by_region.csv and best_selling_products.csv")

```

```

import yfinance as yf
import matplotlib.pyplot as plt
tickers = ["RELIANCE.NS", "TCS.NS", "INFY.NS"]
data = yf.download(tickers, start="2022-10-01", end="2023-10-01",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['RELIANCE.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="Reliance Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="Reliance Industries - Daily Returns",
color='orange')
plt.tight_layout()
plt.show()
reliance_data.to_csv('reliance_stock_data.csv')

```

```

tickers = ["HDFCBANK.NS", "ICICI.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['HDFCBANK.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="HDFC Industries - Closing Price")
plt.subplot(2, 1, 2)

```

```

reliance_data['Daily Return'].plot(title="HDFCIndustries - Daily Returns",
color='red') plt.tight_layout()
plt.show()
reliance_data.to_csv('hdfc_stock_data.csv')
print("\nhdfc stock data saved to 'hdfc_stock_data.csv'.")

```

```

tickers = ["HDFCBANK.NS", "ICICIBANK.NS",
"KOTAKBANK.NS"] data = yf.download(tickers,
start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['ICICIBANK.NS']
print("\nSummary statistics for ICICI Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="ICICI Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="ICICI Industries - Daily Returns",
color='BLACK') plt.tight_layout()

```

```

plt.show()
reliance_data.to_csv('icici_stock_data.csv')
print("\nicici stock data saved to 'icici_stock_data.csv'.")

```

```

tickers = ["HDFCBANK.NS", "ICICI.NS",
"KOTAKBANK.NS"] data = yf.download(tickers,
start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['KOTAKBANK.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] =
reliance_data['Close'].pct_change() print("\n")
plt.figure(figsize=(12, 6))

```

```
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="KOTAK Industries -
Closing Price") plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="kotak Industries - Daily Returns",
color='red') plt.tight_layout()
plt.show()
reliance_data.to_csv('kotak_stock_data.csv')
print("\nkotak stock data saved to 'kotak_stock_data.csv'.")
```


Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:

4/3/25 LAB-2

apsara
Date: _____

Write python code, consider filename as 'housing.csv'

- 1) To load csv file into the dataframe
import pandas as pd
df = pd.read_csv('file_path');
- 2) To display information of all columns
Code: print("Dataset Information:")
print(df.info())

Output: # Column Non-Null Count Dtype
0 longitude 20640 non-null float64
- 3) To display statistical information of all numerical
Code: print("Statistical Information:")
print(df.describe())

Output: longitude latitude housing_median_age total_rooms
mean -119.569 35.63 28.63 2635.76
total_bedrooms population households median_income
537.87 1425.47 499.53 3.87
Median_house_value
206855.816

AP
4/3/25

Page No.: _____

iv) To display the count of unique labels for 'Ocean Proximity' column

→ Code: `print("Unique Labels Count for 'Ocean Proximity'")
print(df['ocean-proximity'].value_counts())`

Output: Unique Labels Count for 'Ocean Proximity' column

Little Ocean 3134

Inland 6551

Near Ocean 2458

g) To display which attributes in a dataset having missing values count greater than 0

→ Code: `missing_values = df.isnull().sum()
columns_with_missing = missing_values[missing_values > 0]
print(columns_with_missing)`

Output: Attributes with Missing Values:

total_bedrooms 207

dtype: int64

For Diabetes and Adult Income

1) Which columns in dataset had missing values? How did you handle them?

2) Which categorical column did you identify in the dataset? How did you encode them?

3) What is the difference between Min Max Scaling and Standardization? When would you use one over the other?

1) Missing values are present in numerical columns if present which are replaced by mean of the respective column

2) No categorical columns, hence no encoding

3) Min Max Scaling transform data to a fixed range [0, 1] using

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

It is used when dataset doesn't follow a normal distribution, have different ranges and need to be bound

Standardization transforms data to have zero mean and unit variance

$$x' = \frac{x - \mu}{\sigma}$$

It is used when dataset follows a

Page No.

Gaussian distribution, many ML algorithms assume normality.

Diabetes.csv

import pandas as pd

import numpy as np

from sklearn.preprocessing import MinMaxScaler, StandardScaler

import seaborn as sns

import matplotlib.pyplot as plt

import sklearn

file_path = "diabetes.csv"

df = pd.read_csv(file_path)

df_num = df.select_dtypes(include=['numeric']).copy()

inputs = df.select_dtypes(include=['numeric']).copy()

df_num.loc[:, :] = inputs.fit_transform(df_num)

df[df_num.columns] = df_num

Q1 = df_num.quantile(0.25)

Q3 = df_num.quantile(0.75)

IQR = Q3 - Q1

df = df[~((df_num < (Q1 - 1.5 * IQR)) | (df_num > (Q3 + 1.5 * IQR)))].copy()

min_max_scaler = MinMaxScaler()

df_min_max = pd.DataFrame(min_max_scaler.fit_transform(df_num), columns=df_num.columns)

std_scaler = StandardScaler()

df_std = pd.DataFrame(std_scaler.fit_transform(df_min_max), columns=df_min_max.columns)

print(df_min_max.head())

print(df_std.head())

Adult Income Dataset

1) Missing values are represented by "?" which we replace with NaN

for numerical columns - replace with mean

for categorical columns - replace with mode

2) workclass, educat, marital-status, occupat, relationship, race, sex, native-country, income

Encoding method: Label Encoding to convert categorical to numerical formats

Code:

```
from google.colab import files
diabetes=files.upload()

from google.colab import files
adult_income=files.upload()

df1=pd.read_csv("Dataset of Diabetes .csv")
df1.head()

df2=pd.read_csv("adult.csv")
df2.head()

df1.info()
df2.info()
df1.describe()
df2.describe()

missing_values1 = df1.isnull().sum()
print(missing_values1)
missing_values2 = df2.isnull().sum()
print(missing_values2)

df1['Gender'] = df1['Gender'].replace('f', 'F')
ordinal_encoder = OrdinalEncoder(categories=[["F", "M"]])
df1["Gender_Encoded"] =
ordinal_encoder.fit_transform(df1[["Gender"]]) onehot_encoder =
OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df1[["CLASS"]]) encoded_array
= encoded_data.toarray()
encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["CLASS"])) df_encoded = pd.concat([df1,
encoded_df], axis=1)
df1 = pd.concat([df1, encoded_df], axis=1)
df1.drop("CLASS", axis=1, inplace=True)
df1.drop("Gender", axis=1, inplace=True)
print(df2.head())
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
df_copy2 = df2
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
df_copy2["Gender_Encoded"] =
ordinal_encoder.fit_transform(df_copy2[["gender"]])
print(df_copy2[["gender", "Gender_Encoded"]])

onehot_encoder = OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df2[["occupation", "workclass", "education",
"marital status", "relationship", "race", "native-country", "income"]])
encoded_array = encoded_data.toarray()
encoded_df =
pd.DataFrame(encoded_array,
```



```

columns=onehot_encoder.get_feature_names_out(["occupation","workclass","education",
", "marital status","relationship","race","native-country","income"]))
df_encoded = pd.concat([df_copy2, encoded_df], axis=1)

df_encoded.drop("gender", axis=1, inplace=True)
df_encoded.drop("occupation", axis=1, inplace=True)
df_encoded.drop("workclass", axis=1, inplace=True)
df_encoded.drop("education", axis=1, inplace=True)
df_encoded.drop("marital-status", axis=1, inplace=True)
df_encoded.drop("relationship", axis=1, inplace=True)
df_encoded.drop("race", axis=1, inplace=True)
df_encoded.drop("native-country", axis=1, inplace=True)
df_encoded.drop("income", axis=1, inplace=True)
print(df_encoded.head())

normalizer = MinMaxScaler()
df_encoded[["fnlwgt","educational-num","capital-gain","capital-loss","hours-per-week"]] =
normalizer.fit_transform(df_encoded[["fnlwgt","educational-num","capital-gain","capital-
loss","hours-per week"]
])
df_encoded.head()
normalizer = MinMaxScaler()
df1[["No_Pation","AGE","Urea","Cr" ,"HbA1c" ,
"Chol","TG","HDL","LDL","VLDL","BMI"]] =
normalizer.fit_transform(df1[["No_Pation","AGE","Urea","Cr" ,"HbA1c" ,
"Chol","TG","HDL","LDL","VLDL","BMI"]])
df1.head()

```

Program 3

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot:

The image shows two pages of handwritten notes. The left page, titled 'LAB-4', describes a linear regression problem using matrix algebra. It lists data points (x, y) and shows the calculation of the regression line $y = -0.5 + 2.2x$. The right page, titled 'Linear Regression', provides a Python code snippet to solve the same problem using the sklearn library. The code imports pandas and sklearn, reads a CSV file, fits a linear regression model, and prints the predicted income for the year 2020.

LAB-4

Linear Regression using Matrix approach

x_i (week)	y_i (sales in k)
1	2
2	4
3	5
4	9

$y = a_0 + a_1x$

$X = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix}$ $\beta = \begin{bmatrix} a_0 \\ a_1 \end{bmatrix}$

$\beta = (X^T X)^{-1} X^T Y$

$X^T X = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}$ $(X^T X)^{-1} = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix}$

$(X^T X)^{-1} X^T = \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix}$

$((X^T X)^{-1} X^T) Y = \begin{bmatrix} -0.5 \\ 2.2 \end{bmatrix}$

$y = -0.5 + 2.2x$

Linear Regression

1) Predict Canada's per capita income in year 2020
Use the data file Canada-per-capita-income.csv file. If required, apply the regression model and predict per capita income for Canadian citizen in 2020

→

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

data = pd.read_csv('Canada-per-capita-income.csv')
x = data['year']
y = data['per capita income (US$)']
model = LinearRegression()
model.fit(x, y)
c = model.coef_
m = model.intercept_

print(f"slope: {c[0]}, 2f3")
print(f"intercept: {m}, 2f3")
year_2020 = np.array([2020])
income_2020 = model.predict(year_2020)
print(f"Predicted capita income 2020: {income_2020}, 2f3")
```

apsara
Date: _____

```

plt.scatter(x, y, color='blue', label='Actual Data')
plt.plot(x, model.predict(x), color='red', label='Regression Line')
plt.xlabel('x-axis')
plt.ylabel('Per capita income (US$)')
plt.show()
print(f"MSE: {mse} mse: {mse}")

Output: slope = 828.47      2nd m = 9318.44
      intercept = -1632210.76 c = 262891
Predicted per capita income in 2020: 94288.61
MSE: 1546939.06

```

Ex:

- Multiple Linear Regression

1) Consider data file 1000_companies.csv the file contains profit status for a firm such as R&D, Administration, Marketing Spend and State. Based on these 4 factors build Multiple Linear Regression Model to predict profit.

R&D = 91634.74
Administration = 11931.24
Marketing Spend = 515841.3
State : Florida

apsara
Date: _____

```

from sklearn.preprocessing import LabelEncoder
data = pd.read_csv("1000_companies.csv")
data['state'] = data['state'].astype('category')
label_encoder = LabelEncoder()
data['state'] = label_encoder.fit_transform(
    data['state'])
x = data.iloc[:, 1:5]
y = data.iloc[:, 5]
model = LinearRegression()
model.fit(x, y)
test_data = np.array([[91634.74, 11931.24, 515841.3, 11931.24], label_encoder.transform(['Florida'])])
predicted_profit = model.predict(test_data)
print(f"Predicted profit: {predicted_profit[0]}")
print(f"slope: {model.coef_[0]}")
print(f"Intercept (c): {model.intercept_[0]}")
y_pred = model.predict(x)
mse = mean_squared_error(y, y_pred)
print(f"MSE: {mse} mse: {mse}")

```

Output: Predicted profit: 9253.01
 slope: 0.55
 Intercept = -70214.74
 MSE: 92128865.28

③ Candidate 1 salary: 4773289
 Candidate 2 salary: 86424.67
 Intercept 14710.23
 slope = 2633.05

11/3/25

Code:

```
from google.colab import files
per_capita_income=files.upload()

from google.colab import files
salary=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import
train_test_split from sklearn.impute
import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder,
OneHotEncoder from sklearn.preprocessing import
StandardScaler, MinMaxScaler from scipy import stats
from sklearn import linear_model

df1=pd.read_csv("canada_per_capita_inc
ome.csv") df1.head()

df2=pd.read_csv("salary.csv")
df2.head()
df2.YearsExperience.median()
df2.YearsExperience =
df2.YearsExperience.fillna(df2.YearsExperience.median()) df2

plt.xlabel("year")
plt.ylabel("per capita income (US$)")
plt.scatter(df1.year, df1['per capita income (US$)'])

plt.xlabel("YearsExperience")
plt.ylabel("Salary")
plt.scatter(df2.YearsExperience, df2.Salary)

reg1 = linear_model.LinearRegression()
reg1.intercept_
reg1.predict([[2020]])

reg2 = linear_model.LinearRegression()
reg2.fit(df2.drop('Salary', axis='columns'),
df2['Salary']) reg2.coef_
reg2.intercept_
reg2.predict([[12]])

from google.colab import files
hiring=files.upload()

from google.colab import files
companies=files.upload()

df3=pd.read_csv("hiring.csv")
```

```

df3.head()

df4=pd.read_csv("1000_Companies.csv")
df4.head()

df3.isnull().sum()
df4.isnull().sum()

df3_copy = df3.copy()
experience_mapping = {'two': 2, 'three': 3, 'five': 5, 'seven': 7, 'ten': 10,
'eleven': 11} df3_copy['experience'] =
df3_copy['experience'].map(experience_mapping)
median_experience = df3_copy['experience'].median()
df3_copy['experience'] = df3_copy['experience'].fillna(median_experience)
df3_copy
df3_copy['test_score(out of 10)'] = df3_copy['test_score(out of
10)'].fillna(df3_copy['test_score(out of 10)'].mean())
reg3 = linear_model.LinearRegression()
reg3.fit(df3_copy.drop('salary($)', axis='columns'),
df3_copy['salary($)']) reg3.coef_
reg3.intercept_
reg3.predict([[2,9,6]])
reg3.predict([[12,10,10]])
ohe = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
state_encoded = ohe.fit_transform(df4[['State']])
state_encoded_df = pd.DataFrame(state_encoded, columns=ohe.get_feature_names_out(['State']))

df4 = pd.concat([df4, state_encoded_df], axis=1).drop(columns=['State'])
print(df4)
reg4 = linear_model.LinearRegression()
reg4.fit(df4.drop('Profit',axis='columns'),df4.Profit)
print(reg4.coef_)
print(reg4.intercept_)
reg4.predict([[91694.48, 515841.3, 11931.24,0,1,0]])

```


Program 4

Build Logistic Regression Model for a given dataset

Screenshot:

apsara
Date: _____

LAB - 3

i) Given $a_0 = -5$ $a_1 = 0.8$

ii) Logistic regression equation

$$p(x) = \frac{1}{1 + e^{-(a_0 + a_1 x)}} = \frac{1}{1 + e^{-(-5 + 0.8x)}}$$

ii) Calculate probability that a student who studies for 7 hrs will pass

→ $x = 7$ $p(x) = \frac{1}{1 + e^{-(-5 + 0.8(7))}}$

$$= 0.6457$$

iii) Determine the predicted class (PIF) for this student based on threshold of 0.5

→ $P(x) = 0.6457$

$$P(x) \geq 0.5 \quad y = \begin{cases} 1 & \text{if } P(x) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Thus $y = 1$ (Pass)

2) Consider $z = [2, 1, 0]$ for three classes

Apply Softmax function to find probability values of 3 classes

→ $\text{Softmax}(z_k) = \frac{e^{z_k}}{\sum_{j=1}^K e^{z_j}}$

$$\text{Softmax}(z_1) = \frac{e^2}{e^2 + e^1 + e^0} = 0.665$$

Page No.: _____

$$\text{Softmax}(z_2) = \frac{e^0}{e^0 + e^0 + e^0} = 0.244$$

$$\text{Softmax}(z_3) = \frac{e^0}{e^0 + e^0 + e^0} = 0.091$$

Probabilities of the 3 classes are approximately 66.5%, 24.4% & 9.1%.

- Logistic Regression - Binary
 - import pandas as pd
 - from matplotlib import pyplot as plt
 - df = pd.read_csv('insurance_data.csv')
 - df.head()
 - plt.scatter(df['age'], df['bought'], marker='x', color='red')
 - from sklearn.model_selection import train_test_split
 - X_train, X_test, y_train, y_test = train_test_split(df[['age']], df['bought'], train_size=0.9, random_state=10)
 - X_train.shape
 - X_test
 - from sklearn.linear_model import LogisticRegression
 - model = LogisticRegression()

```
model.fit(X_train, y_train)
X_test
y_test
y_predicted = model.predict(X_test)
model.score(X_test, y_test)
model.predict(X_test)
y_predicted = model.predict([[60]])
y_predicted
model.coef
model.intercept
import math
def sigmoid(z):
    return 1/(1+math.exp(-z))
def prediction_func(age):
    z = 0.127 * age - 4.973
    y = sigmoid(z)
    return y
age = 35
prediction_func(age)
```

Output:
0.35 is less than 0.5 which means person with age 35 will not buy insurance

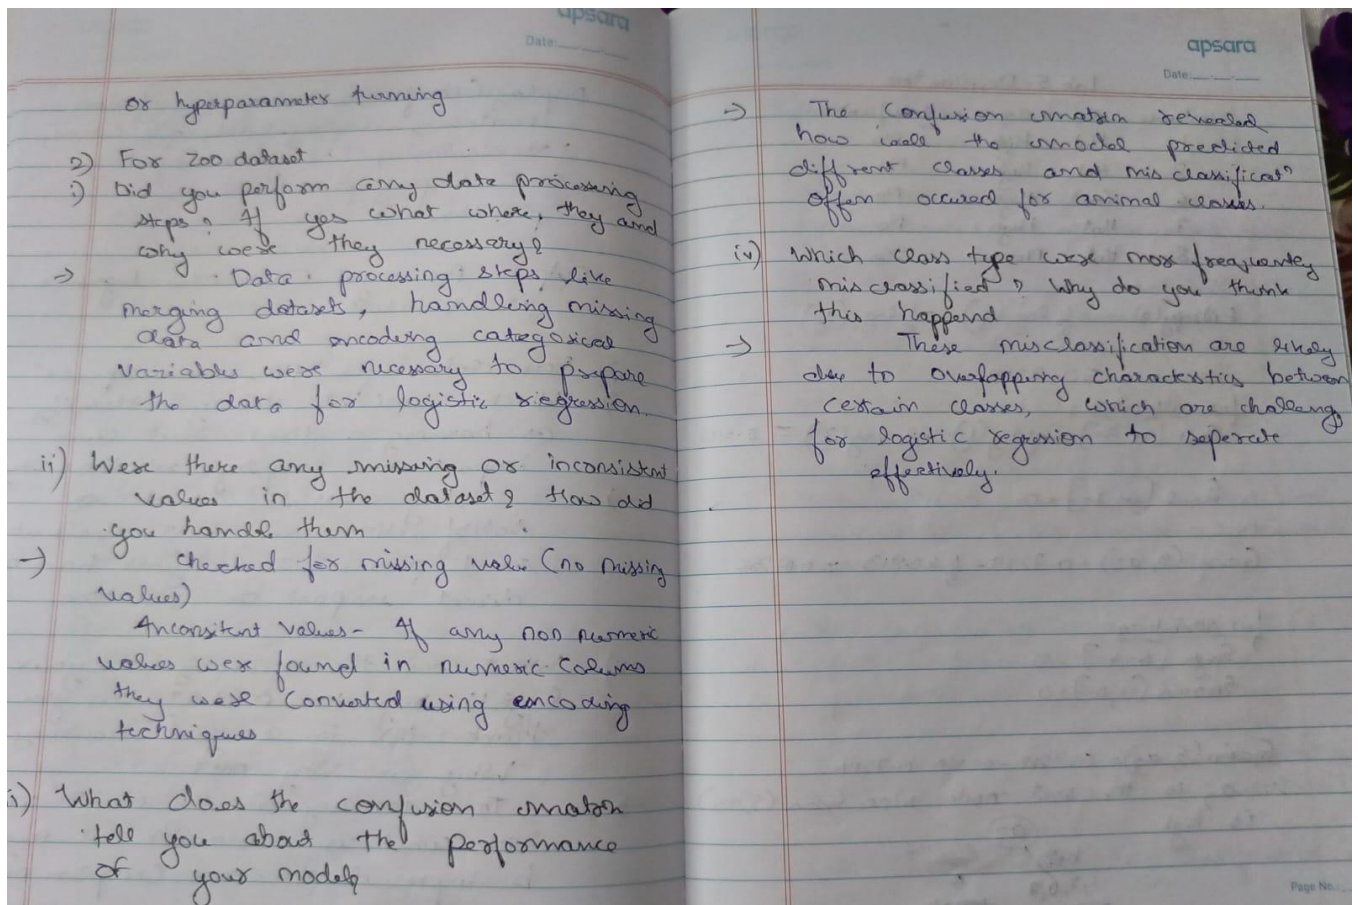
- Logistic Regression - Multiclass
 - import pandas as pd
 - from sklearn.datasets import load_iris
 - from sklearn.linear_model import LogisticRegression
 - from sklearn.metrics import accuracy_score
 - from sklearn import metrics
 - import matplotlib.pyplot as plt
 - iris = pd.read_csv('iris.csv')
 - iris.head()
 - X = iris.drop('species', axis='columns')
 - y = iris['species']
 - X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
 - model = LogisticRegression(multi_class='multinomial')
 - model.fit(X_train, y_train)
 - y_pred = model.predict(X_test)
 - accuracy = accuracy_score(y_test, y_pred)
 - print('accuracy of Multinomial Regression {accuracy: .2f}')
 - Confusion Matrix = metrics.confusion_matrix(y_test, y_pred)
 - cm_display = metrics.ConfusionMatrix

```
Display Confusion Matrix = confusion_matrix, display_labels = ['setosa', 'versicolour', 'virginica'])
cm_display.plot()
plt.show()
```

Output: Accuracy of Multinomial Regression is set to 1.00

i) For dataset file "HR_comma_sep.csv"
ii) Which variable did you identify as having a direct and clear impact on employee retention? Why?
Key variables such as satisfaction level, last evaluation and promoted. Last 5 years had a direct impact on retention

ii) What was the accuracy of your logistic regression model? Do you think this is a good accuracy? Why or why not?
The accuracy of the model was 29.1% which is reasonable but might be improved with more complex model



Code:

```
from google.colab import files
hr=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder,
OneHotEncoder from sklearn.preprocessing import
StandardScaler, MinMaxScaler
from scipy import stats
from sklearn import linear_model
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

df1=pd.read_csv("HR_comma_sep.csv")
df1.head()
df1.isnull().sum()
plt.figure(figsize=(12, 6))
```

```

sns.barplot(x='Department', y='left', data=df1)
plt.title('Employee Retention Rate by Department')
plt.xlabel('Department')
plt.ylabel('Proportion of Employees Left')
plt.xticks(rotation=45, ha='right')
plt.show()

ohe = OneHotEncoder(handle_unknown='ignore',
sparse_output=False) department_encoded =
ohe.fit_transform(df1[['Department']])
department_encoded_df = pd.DataFrame(department_encoded,
columns=ohe.get_feature_names_out(['Department']))
df1 = pd.concat([df1, department_encoded_df], axis=1)
df1 = df1.drop('Department', axis=1)
ordinal_encoder = OrdinalEncoder(categories=[['low', 'medium', 'high']],
dtype=np.int64) salary_encoded =
ordinal_encoder.fit_transform(df1[['salary']])
df1['salary_encoded'] = salary_encoded
df1 = df1.drop('salary', axis=1)
df1.head()

```

```

correlation_matrix = df1.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f") plt.title('Correlation Matrix of Features')
plt.show()
plt.figure(figsize=(8, 6))
sns.barplot(x='salary_encoded', y='left', data=df1)
plt.title('Impact of Employee Salary on Retention')
plt.xlabel('Salary Level (Encoded)')
plt.ylabel('Proportion of Employees Left')
plt.show()

```

```

df_copy = df1[['number_project', 'average_monthly_hours', 'time_spend_company',
'left', 'salary_encoded', 'satisfaction_level', 'Work_accident']]
df_copy.head()
X = df_copy.drop('left', axis=1)
y = df_copy['left']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the Logistic Regression model: {accuracy}")

```

```

from google.colab import files
zoodata=files.upload()
zootype=files.upload()

```

```

zoo_data = pd.read_csv('zoo-data.csv')
zoo_class = pd.read_csv('zoo-class-type.csv')

```

```

merged_data = pd.merge(zoo_data, zoo_class, left_on='class_type',
right_on='Class_Number') merged_data = merged_data.drop(['Animal_Names',
'Number_Of_Animal_Species_In_Class',
'Class_Number','class_type','animal_name'], axis=1)
X = merged_data.drop('Class_Type', axis=1)
y = merged_data['Class_Type']
print(merged_data.head())
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=np.unique(y_test)) disp.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix")
plt.show()

```

Program 5

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:

Lab 5 - Decision tree

instance	a ₁	a ₂	classification
1	Hot	High	No
2	Hot	High	No
6	Cool	High	No
7	Hot	High	No
8	Hot	Normal	Yes

$$\text{Entropy}(S) = -\frac{4}{9} \log_2 \left(\frac{4}{9} \right) - \frac{1}{9} \log_2 \left(\frac{1}{9} \right)$$

$$= 0.7219$$

for attribute a₁:

$$S_{\text{hot}} [1, 2, 6, 7] = -\frac{4}{4} \log_2 \left(\frac{4}{4} \right) - \frac{3}{4} \log_2 \left(\frac{3}{4} \right) = 0.8113$$

$$S_{\text{cool}} [0, 1] = 0$$

$$\text{Gain}(S, a_1) = 0.7219 - \frac{4}{9} \times 0.8113 = 0.4986$$

for attribute a₂:

$$S_{\text{high}} [0, 5, 6, 7] = 0$$

$$S_{\text{normal}} [1, 8] = 0$$

$$\text{Gain}(S, a_2) = 0.7219 - 0 = 0.7219$$

∴ a₂ is the root node since Gain(S, a₂) is high

```

graph TD
    A2((a2)) -- High --> B[1, 2, 6, 7]
    A2 -- Normal --> C[8]
    B --> D[No]
    C --> E[Yes]
        
```

Python code for Decision tree

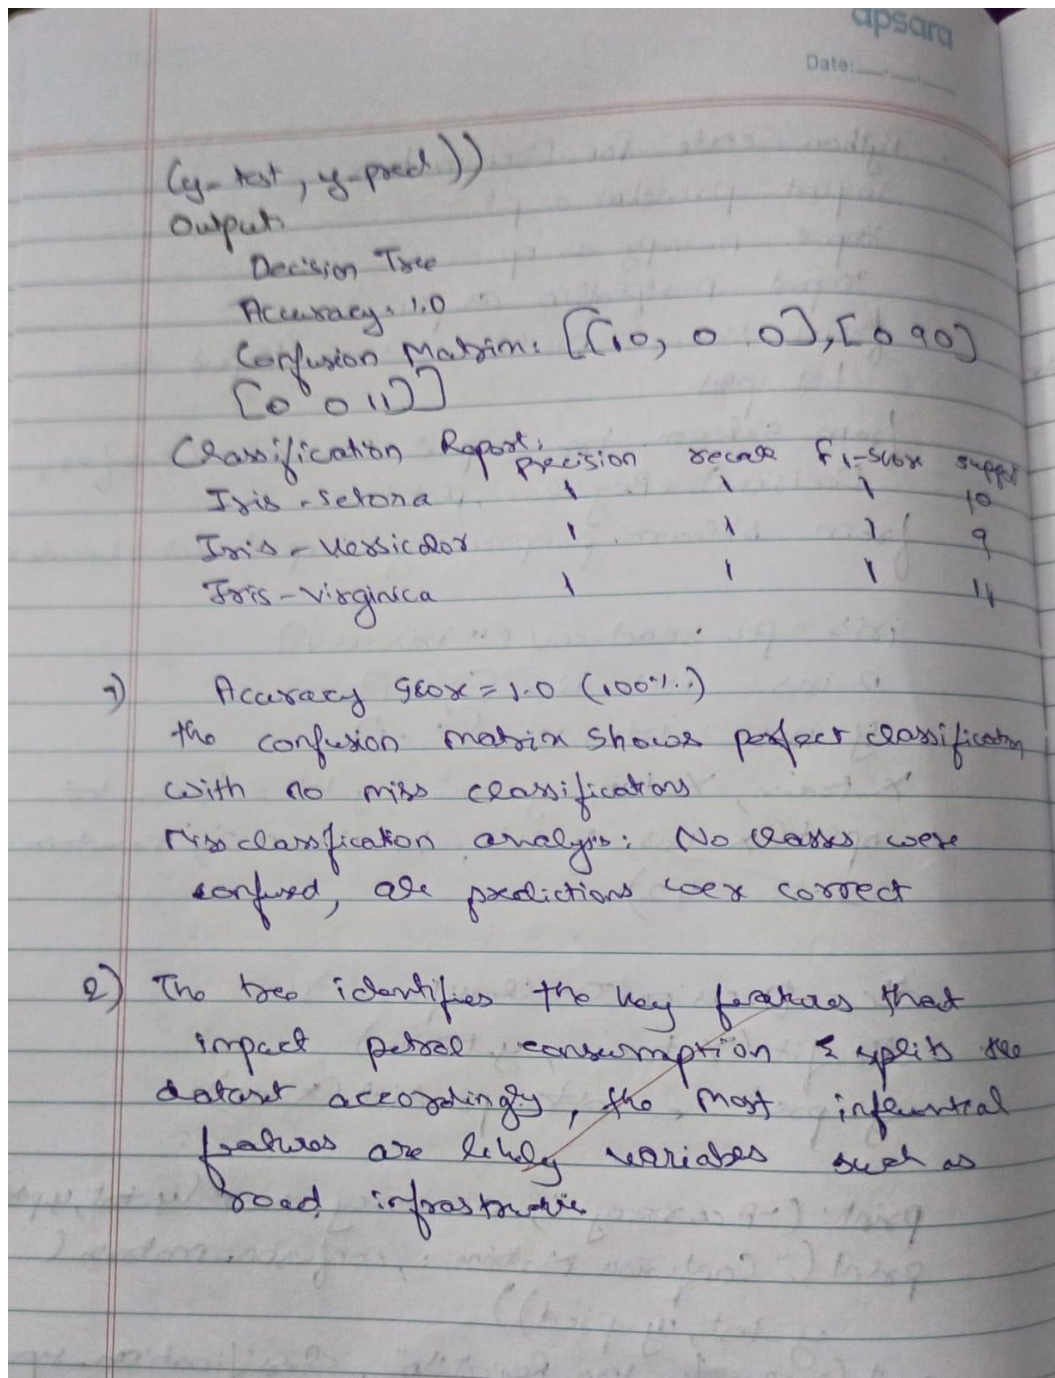
```

import pandas as pd
import numpy as np
import matplotlib as mp
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.preprocessing import LabelEncoder

iris = pd.read_csv("iris.csv")
X = iris.iloc[:, :-1]
y = iris.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

dtc = DecisionTreeClassifier()
dtc.fit(X_train, y_train)
y_pred = dtc.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Confusion Matrix:", confusion_matrix(y_test, y_pred))
print("Classification Report:", classification_report(y_test, y_pred))
        
```



Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris.csv")
df1.head()
```

```
df1.isnull().sum()
```

```
X = df1.drop('species', axis=1)
y = df1['species']
```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred))
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=X.columns,
class_names=y.unique()) plt.show()

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=clf.classes_) cmap = plt.cm.get_cmap('PuBuGn')
disp.plot(cmap=cmap)
plt.show()

drug=files.upload()
df2=pd.read_csv("drug.csv")
df2.head()
df2.isnull().sum()

label_encoders = {}
for column in df2.columns:
    le = LabelEncoder()
    df2[column] = le.fit_transform(df2[column])
    label_encoders[column] = le
X = df2.drop('Drug', axis=1)
y = df2['Drug']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred))
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=X.columns, class_names=[str(c) for c in
y.unique()]) plt.show()

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
cmap = plt.cm.Blues
disp.plot(cmap=cmap)
plt.show()

pc=files.upload()
df3=pd.read_csv("petrol_consumption.csv")
df3.head()
df3.isnull().sum()
X = df3.drop('Petrol_Consumption', axis=1)
y = df3['Petrol_Consumption']

```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) regressor = DecisionTreeRegressor(random_state=42)
regressor.fit(X_train, y_train)
y_pred =
regressor.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse:.2f}')
print(f'Root Mean Squared Error:
{rmse:.2f}') print(f'Mean Absolute Error:
{mae:.2f}') print(f'R-squared: {r2:.2f}')
plt.figure(figsize=(30, 30))
plot_tree(regressor, filled=True, feature_names=X.columns,
fontsize=10) plt.show()

```

Program 6

Build KNN Classification model for a given dataset.

Screenshot :

apsara
Date: _____

LAB-6 KNN

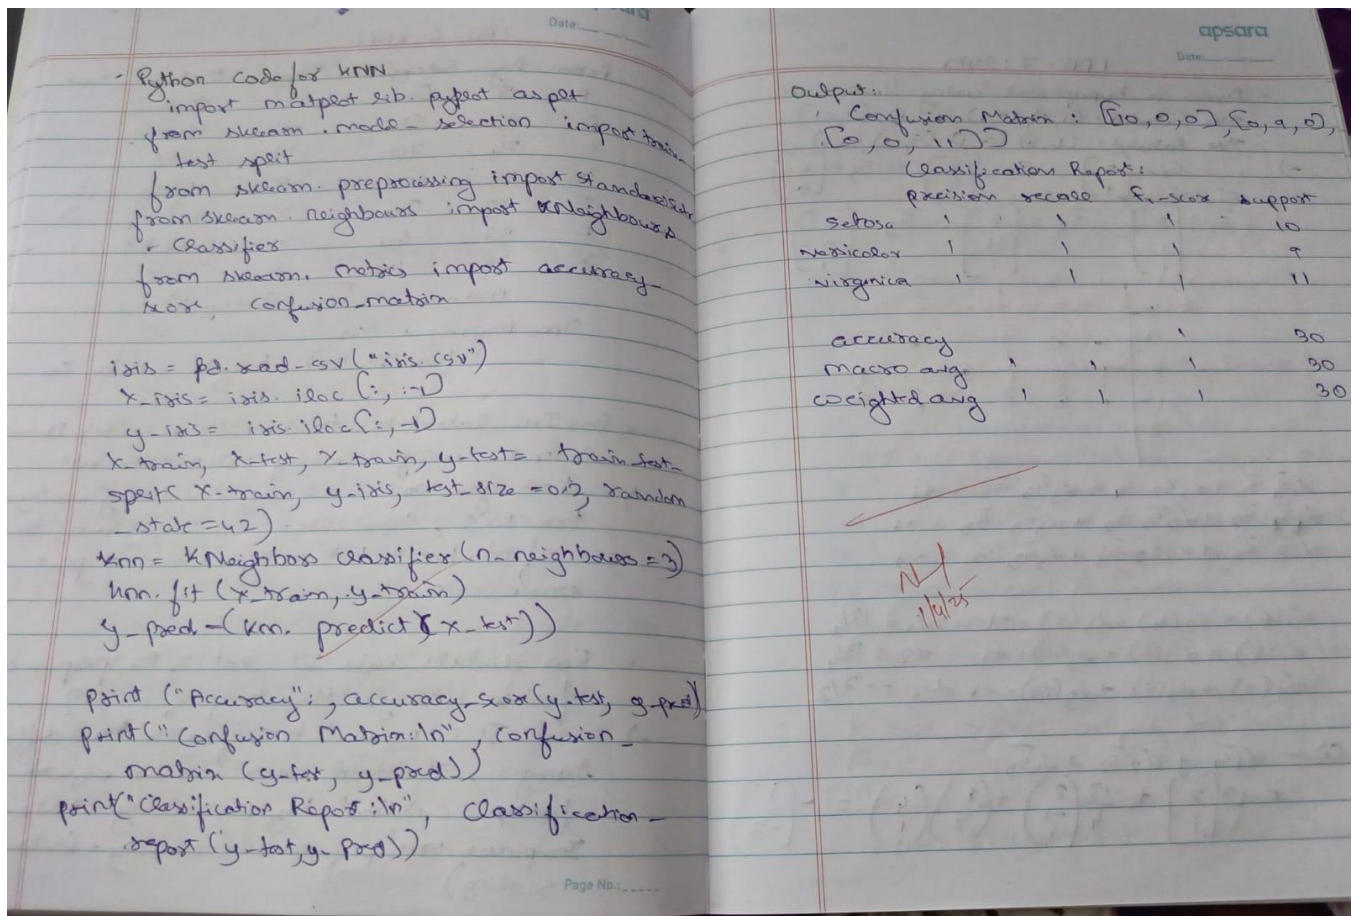
Person	Age	Salary(k)	Target	Distance	Rank
A	18	50	N	52.8	
B	23	55	N	46.57	
C	24	70	N	31.95	
D	41	60	Y	40.44	2
E	43	70	Y	31.04	3
F	38	60	Y	60.07	1
X	35	100	Y		

$\text{Distance} = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$

Since Majority is Yes
for (35, 100) target will be Yes

- For Iris dataset how to choose the k value? Demonstrate using accuracy rate & error rate
- K=3 gives 100% accuracy but generally take K=5, error rate = 1 - accuracy = 0, to find best k, test multiple values & plot error rate
- For diabetes data set what is the purpose of feature scaling? How to perform it?
- It is needed bcz features have different ranges, standard scaler ensures equal feature contribution by normalizing data. Improves KNN performance (accuracy = 69.84% after scaling)

Page No.: _____



Code:

```

from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (2).csv")
df1.head()
df1.isnull().sum()
X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}, Error Rate for k={k}: {1-accuracy}")
    if accuracy > best_accuracy: best_accuracy = accuracy best_k = k
print(f"Best k value: {best_k}")
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, y_pred))

```

```

print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
print(cm)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=knn.classes_, yticklabels=knn.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

diabetes=files.upload()
df2=pd.read_csv("diabetes.csv")
df2.head()
df2.isnull().sum()
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df2.drop('Outcome', axis=1))
X_train, X_test, y_train, y_test = train_test_split(X_scaled, df2['Outcome'], test_size=0.2,
random_state=42) best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
print(f"Best k value: {best_k}")
knn = KNeighborsClassifier(n_neighbors=best_k)
knn.fit(X_train, y_train) y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted") plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

heart=files.upload()
df3=pd.read_csv("heart.csv")
df3.head()

```

```

df3.isnull().sum()
X = df3.drop('target', axis=1)
y = df3['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}, Error Rate for k={k}: {1-
accuracy}") if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
print(f"Best k value: {best_k}")
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
print(cm)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
xticklabels=knn.classes_, yticklabels=knn.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```


Program7

Build Support vector machine model for a given dataset

Screenshot:

LAB - 7 : SVM

Draw a hyperplane using linear SVM
 The labeled: $(1, 1), (2, 1), (3, -1), (3, -1)$
 The labeled: $(1, 0), (2, 1), (3, -1), (3, 0)$

→ $S_1 = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$ $S_2 = \begin{bmatrix} 2 \\ -1 \end{bmatrix}$ $S_3 = \begin{bmatrix} 4 \\ 0 \end{bmatrix}$

$\bar{S}_1 = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$ $\bar{S}_2 = \begin{bmatrix} 2 \\ -1 \end{bmatrix}$ $\bar{S}_3 = \begin{bmatrix} 4 \\ 0 \end{bmatrix}$

$\alpha_1 \bar{S}_1 \cdot \bar{S}_1 + \alpha_2 \bar{S}_2 \cdot \bar{S}_2 + \alpha_3 \bar{S}_3 \cdot \bar{S}_3 = +1$
 $\alpha_1 \bar{S}_1 \cdot \bar{S}_2 + \alpha_2 \bar{S}_2 \cdot \bar{S}_2 + \alpha_3 \bar{S}_3 \cdot \bar{S}_3 = +1$
 $\alpha_1 \bar{S}_1 \cdot \bar{S}_3 + \alpha_2 \bar{S}_2 \cdot \bar{S}_3 + \alpha_3 \bar{S}_3 \cdot \bar{S}_3 = -1$

$\alpha_1(6) + \alpha_2(4) + \alpha_3(9) = +1$ $\alpha_1 = 13/4$
 $\alpha_1(4) + \alpha_2(6) + \alpha_3(9) = +1$ $\alpha_2 = 13/4$
 $\alpha_1(9) + \alpha_2(9) + \alpha_3(12) = -1$ $\alpha_3 = -7/2$

$C = \alpha_1 \bar{S}_1 + \alpha_2 \bar{S}_2 + \alpha_3 \bar{S}_3$
 $= \frac{13}{4} \begin{bmatrix} 2 \\ 1 \end{bmatrix} + \frac{13}{4} \begin{bmatrix} 2 \\ -1 \end{bmatrix} + \left(-\frac{7}{2}\right) \begin{bmatrix} 4 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ -3 \end{bmatrix}$

Code

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix

iris_df = pd.read_csv("iris(2).csv")
X_iris = iris_df.drop("species", axis=1)
y_iris = iris_df["species"]

label_encoder = LabelEncoder()
X_iris_encoded = label_encoder.fit_transform(X_iris)
y_iris_encoded = label_encoder.fit_transform(y_iris)

X_train, X_test_iris, y_train, y_test_iris = train_test_split(X_iris_encoded, y_iris_encoded, test_size=0.2, random_state=42)

svm_linear = SVC(kernel="linear")
```

svm_linear.fit(X_train_iris, y_train_iris)

$y_{pred_linear} = svm_linear.predict(X_test_iris)$

print("Linear kernel accuracy:", accuracy_score(y_test_iris, y_pred_linear))

print("Confusion matrix:", confusion_matrix(y_test_iris, y_pred_linear))

svm_rbf = SVC(kernel="rbf")

svm_rbf.fit(X_train_iris, y_train_iris)

$y_{pred_rbf} = svm_rbf.predict(X_test_iris)$

print("RBF kernel accuracy:", accuracy_score(y_test_iris, y_pred_rbf))

print("Confusion matrix:", confusion_matrix(y_test_iris, y_pred_rbf))

Output

Linear kernel Accuracy: 1.0

Confusion matrix: $\begin{bmatrix} 10 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 0 & 11 \end{bmatrix}$

RBF kernel Accuracy: 1.0

Confusion matrix: $\begin{bmatrix} 10 & 0 & 0 \\ 0 & 9 & 0 \\ 0 & 0 & 11 \end{bmatrix}$

1) IRIS dataset

- Accuracy: 100%
- Best kernel: Both performed equally well
- Reason: IRIS dataset is linearly separable

2) Letter Recognition dataset

- Confusion matrix: most letters classified correctly
- AUC score: 1.0
- Excellent model performance
- Comparison: slightly lower accuracy than IRIS (95%) but handles more complex data

N/P
9/4/25

Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (1).csv")
df1.head()
X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train, y_train)
rbf_y_pred = rbf_svm.predict(X_test)
print("RBF Kernel SVM:")
print("Accuracy:", accuracy_score(y_test, rbf_y_pred))
cm = confusion_matrix(y_test, rbf_y_pred)
sns.heatmap(cm, annot=True, fmt='d',cmap="Blues")
plt.title('Confusion Matrix for RBF Kernel SVM')
plt.xlabel('Predicted')
plt.ylabel('True') plt.show()
print(classification_report(y_test, rbf_y_pred))
linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train, y_train)
linear_y_pred = linear_svm.predict(X_test)
print("\nLinear Kernel SVM:")
print("Accuracy:", accuracy_score(y_test, linear_y_pred))
cm = confusion_matrix(y_test, linear_y_pred)
sns.heatmap(cm, annot=True, fmt='d',cmap="Blues")
plt.title('Confusion Matrix for Linear Kernel SVM')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print(classification_report(y_test, linear_y_pred))
letter=files.upload()
df2=pd.read_csv("letter-recognition.csv")
df2.head()
X = df2.drop('letter', axis=1)
y = df2['letter']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) svm_classifier = SVC(kernel='linear', probability=True)
svm_classifier.fit(X_train, y_train)
y_pred =
svm_classifier.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10,10))
sns.heatmap(cm, annot=True, fmt='d', cmap="Blues")
plt.title('Confusion Matrix for SVM')
plt.xlabel('Predicted')
plt.ylabel('True')
```

```

plt.show()
lb = LabelBinarizer()
lb.fit(y_test)

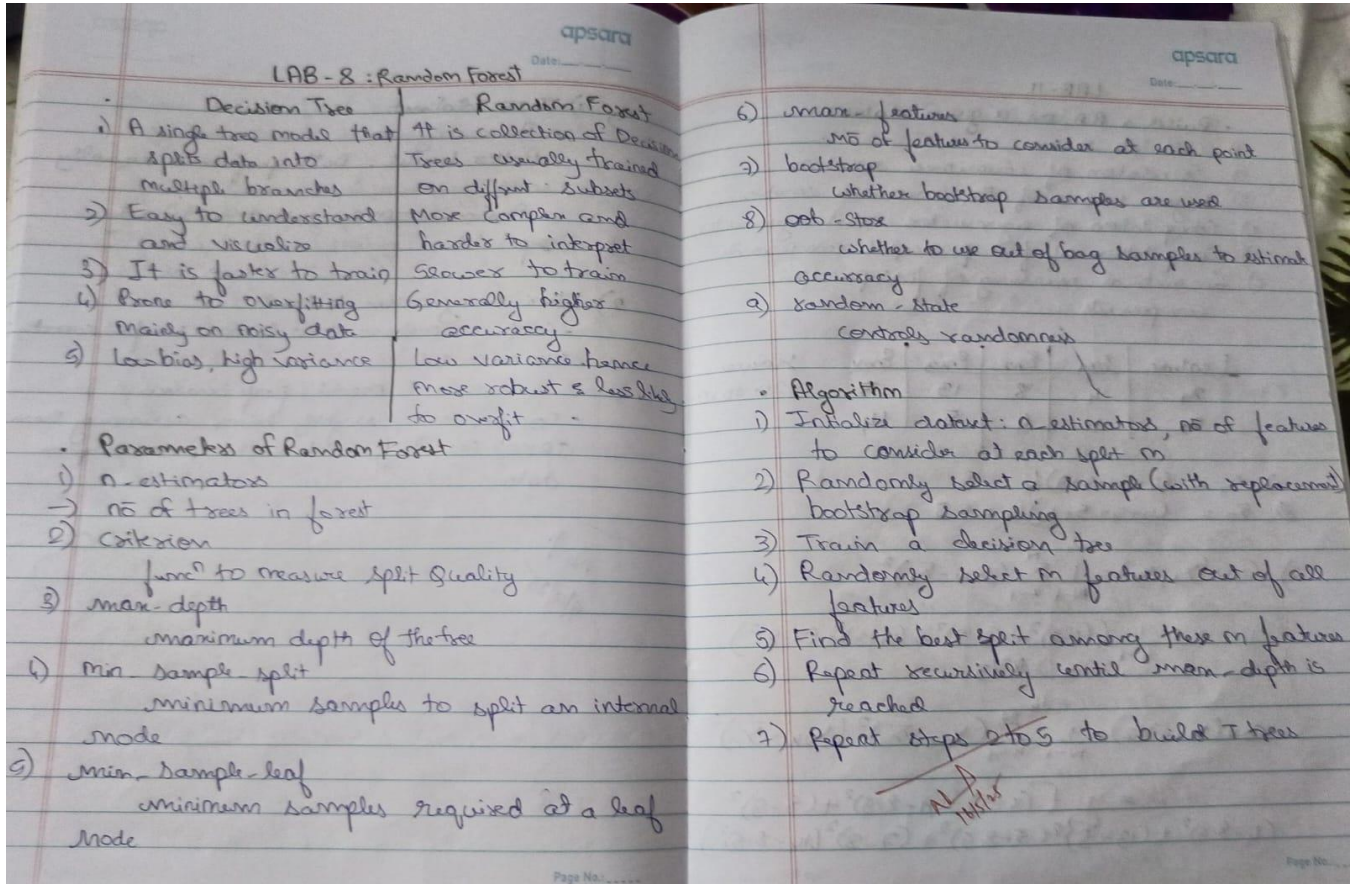
y_test_lb = lb.transform(y_test)
y_pred_prob =
svm_classifier.predict_proba(X_test) fpr = {}
tpr = {}
thresh = {}
roc_auc = dict()
n_class = y_test_lb.shape[1]
for i in range(n_class):
    fpr[i], tpr[i], thresh[i] = roc_curve(y_test_lb[:,i], y_pred_prob[:,i])
    roc_auc[i] = auc(fpr[i], tpr[i])
plt.plot(fpr[0], tpr[0], linestyle='--',color='orange', label='SVM (AUC = %0.2f)' %
roc_auc[0]) plt.title('ROC Curve for Class 0')
plt.xlabel('False Positive
Rate') plt.ylabel('True Positive
rate') plt.legend(loc='best')
plt.show()
print(f"AUC score for class 0: {roc_auc[0]}")

```


Program 8

Implement Random forest ensemble method on a given dataset

Screenshot:



Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (4).csv")
df1.head()
X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=0) rf_classifier = RandomForestClassifier(random_state=0)
rf_classifier.fit(X_train, y_train)
y_pred =
rf_classifier.predict(X_test)
default_accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with default n_estimators: {default_accuracy}")
best_accuracy = 0
best_n_estimators = 0
for n_estimators in range(1, 101):
    rf_classifier = RandomForestClassifier(n_estimators=n_estimators, random_state=0)
    rf_classifier.fit(X_train, y_train)
```

```

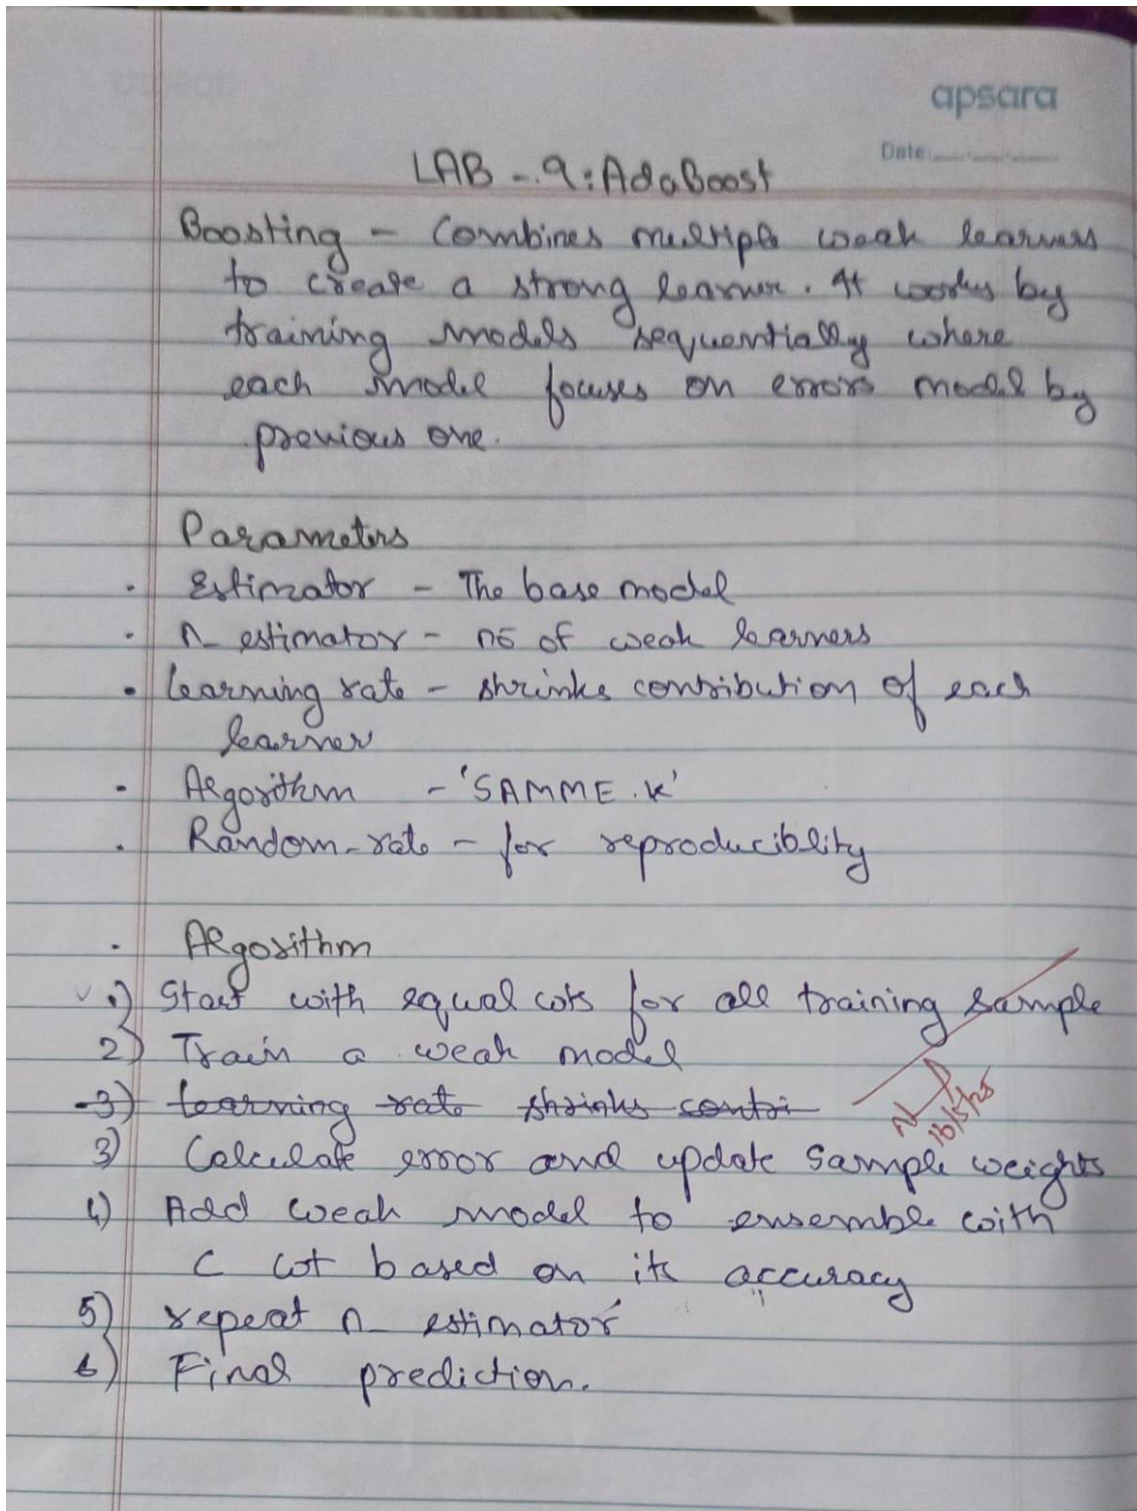
y_pred = rf_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
if accuracy > best_accuracy:best_accuracy
= accuracy best_n_estimators =
n_estimators
print(f"\nBest accuracy: {best_accuracy} achieved with n_estimators =
{best_n_estimators}") cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=np.unique(y_test), yticklabels=np.unique(y_test))
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

```

Program 9

Implement Boosting ensemble method on a given dataset

Screenshot:



Code:

```
from google.colab import files
income=files.upload()
df1=pd.read_csv("income.csv")
df1.head()
X = df1.drop('income_level', axis=1)
y = df1['income_level'] X = pd.get_dummies(X)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42) abc = AdaBoostClassifier(n_estimators=10,
random_state=42)
abc.fit(X_train, y_train)
y_pred = abc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Initial AdaBoost accuracy (10 trees): {accuracy}")
param_grid = {'n_estimators': [50, 100, 150, 200]}
grid_search = GridSearchCV(AdaBoostClassifier(random_state=42), param_grid, cv=5,
scoring='accuracy') grid_search.fit(X_train, y_train)
print(f"Best parameters: {grid_search.best_params_}")
print(f"Best cross-validation score: {grid_search.best_score_}")
best_abc = grid_search.best_estimator_
y_pred_best = best_abc.predict(X_test)
best_accuracy = accuracy_score(y_test,
y_pred_best)
print(f"Accuracy of the best model on the test set: {best_accuracy}")
cm = confusion_matrix(y_test, y_pred_best)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=['<=50K', '>50K'], yticklabels=['<=50K', '>50K'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```


Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file

Screenshot:

apsara
Date: _____

LAB - 10 : K-Means

Cluster the given points into 3 clusters

$A_1(2, 10)$ $A_2(2, 5)$ $A_3(8, 4)$ $A_4(5, 8)$
 $A_5(7, 5)$ $A_6(6, 4)$ $A_7(1, 2)$ $A_8(4, 9)$

Initial cluster $A_1(2, 10)$ $A_4(5, 8)$ $A_7(1, 2)$

→

Given Points	dis from $C_1(2, 10)$	dis from $C_2(5, 8)$	dis from $C_3(1, 2)$	Cluster
$A_1(2, 10)$	0	5	9	C_1
$A_2(2, 5)$	5	6	4	C_3
$A_3(8, 4)$	12	7	4	C_2
$A_4(5, 8)$	5	0	9	C_2
$A_5(7, 5)$	10	5	10	C_2
$A_6(6, 4)$	10	5	9	C_2
$A_7(1, 2)$	9	10	7	C_3
$A_8(4, 9)$	3	2	10	C_2

Center of cluster 1 $C_1(2, 10)$

Center of cluster 2
 $(8+5+7+6+4)/5, (4+8+5+4+9)/5$
 $C_2(6, 6)$

Center of cluster 3
 $(2+1)/2, (5+2)/2$
 $C_3(1.5, 3.5)$

Page No.

Given points	dist from $c_1(2,10)$	dist from $c_2(6,6)$	dist from $c_3(4,3.5)$	Cluster
$P_1(2,10)$	0	8	7	c_1
$P_2(3,5)$	5	5	2	c_3
$P_3(8,4)$	12	4	7	c_2
$P_4(5,8)$	5	3	8	c_2
$P_5(7,5)$	10	2	7	c_3
$P_6(6,4)$	10	2	6	c_2
$P_7(1,2)$	9	9	2	c_3
$P_8(4,9)$	3	5	8	c_1

Center of cluster 1
 $(2+4)/2, (10+9)/2 = (3, 9.5)$

Center of cluster 2
 $(8+5+7+6)/4, (4+8+5+4)/4$
 $= (6.5, 5.25)$

Center of cluster 3
 $(5+1)/2, (5+2)/2$
 $= (1.5, 3.5)$

- Choosing no of clusters
- Elbow method
- Silhouette score
- Domain knowledge

→ Sum of squared errors (SSE)
 $SSE = \sum_{i=1}^n \sum_{k \in C_i} \|x_i - u_k\|^2$

- Algorithm**
- 1) Select the number k to decide no of clusters
 - 2) Select random points or centroids
 - 3) Assign each data point to their closest centroid, which will form the predefined k clusters
 - 4) Calculate the variance and place a new centroid of each cluster
 - 5) Repeat 3rd step which means reassign each datapoint to the new closest centroid of each cluster
 - 6) If any reassignment occurs then go to step 4 else go to Finish
 - 7) The model is ready

- Key parameters**
- $n_clusters$ - no of clusters to form
 - $init$ - initialization method
 - n_init - no of initialization
 - max_iter - Max iteration per run
 - $random_state$ - control randomness
 - tol - convergence threshold
 - $algo$ - algorithm to use

Code:

```
from google.colab import files
iris=files.upload()
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy import stats
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
```

```
df1=pd.read_csv("iris (4).csv")
df1.head()
df = df1.drop(['sepal_length','sepal_width','species'],axis=1)
scaler = StandardScaler()
scaled_df = scaler.fit_transform(df) wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
    random_state=0) kmeans.fit(scaled_df)
```

```
wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
kmeans = KMeans(n_clusters=3, init='k-means++', max_iter=300, n_init=10,
random_state=0) pred_y = kmeans.fit_predict(scaled_df)
df['cluster'] = pred_y
plt.scatter(df['petal_length'], df['petal_width'], c=df['cluster'])
plt.title('Clusters of Iris Flowers')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.show()
```


Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA)

method.

Screenshot:

LAB-11

- Build a PCA on a given dataset
- Calculate mean
- Calculation of covariance matrix
- Eigen values of the covariance matrix
- Computation of eigenvectors - unit eigenvectors
- Computation of first principal components
- Geometrical meaning of first principal components

Feature	x_{11}	x_{12}	x_{13}	x_{14}	reduce
x_1	4	8	13	7	20 to 10
x_2	11	4	5	14	

1) $\bar{x}_1 = (4+8+13+7)/4 = 8$
 $\bar{x}_2 = (11+4+5+14)/4 = 8.5$

2) $Cov(x_1, x_1) = \frac{1}{N-1} \sum_{k=1}^N (x_{1k} - \bar{x}_1)^2$
 $= \frac{1}{4-1} ((4-8)^2 + (8-8)^2 + (13-8)^2 + (7-8)^2)$
 $= \frac{14}{3}$

$Cov(x_1, x_2) = \frac{1}{N-1} \sum_{k=1}^N (x_{1k} - \bar{x}_1)(x_{2k} - \bar{x}_2)$
 $= \frac{1}{4-1} ((4-8)(11-8.5) + (8-8)(4-8.5) + (13-8)(5-8.5) + (7-8)(14-8.5))$
 $= \frac{1}{3} ((-4)(2.5) + (0)(-4.5) + (5)(-3.5) + (-1)(5.5))$
 $= \frac{1}{3} (-10 + 0 - 17.5 - 5.5) = \frac{-33}{3} = -11$

$Cov(x_2, x_1) = -11$
 $Cov(x_2, x_2) = \frac{1}{N-1} \sum_{k=1}^N (x_{2k} - \bar{x}_2)^2$
 $= \frac{1}{4-1} ((11-8.5)^2 + (4-8.5)^2 + (5-8.5)^2 + (14-8.5)^2)$
 $= \frac{1}{3} (6.25 + 20.25 + 12.25 + 30.25) = \frac{69}{3} = 23$

Cov matrix = $\begin{bmatrix} Cov(x_1, x_1) & Cov(x_1, x_2) \\ Cov(x_2, x_1) & Cov(x_2, x_2) \end{bmatrix}$
 $= \begin{bmatrix} 14/3 & -11 \\ -11 & 23 \end{bmatrix}$

3) Finding eigen values
 $0 = \det(S - \lambda I) \Rightarrow \begin{vmatrix} 14/3 - \lambda & -11 \\ -11 & 23 - \lambda \end{vmatrix}$
 $(14/3 - \lambda)(23 - \lambda) - 121 = 0$
 $322 - 37\lambda + \lambda^2 - 121 = 0$
 $\lambda^2 - 37\lambda + 201 = 0$
 $\lambda_1 = 30.88, \lambda_2 = 6.15$

4) $u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$ $\begin{bmatrix} 0 \\ 0 \end{bmatrix} = (S - \lambda_1 I)u$
 $\begin{bmatrix} 14/3 - 30.88 & -11 \\ -11 & 23 - 30.88 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$
 $(14/3 - 30.88)u_1 - 11u_2 = 0$
 $-11u_1 + (23 - 30.88)u_2 = 0$
 $\frac{u_1}{11} = \frac{u_2}{14 - \lambda_1}$
 $u_1 = 11t, u_2 = (14 - \lambda_1)t$
 $u_1 = \begin{bmatrix} 11 \\ 14 - \lambda_1 \end{bmatrix}$

$\|u\| = \sqrt{u_1^2 + u_2^2} = \sqrt{11^2 + (14 - 30.88)^2}$
 $= \sqrt{121 + 224.88} = \sqrt{345.88} = 18.6$

$e_1 = \frac{1}{18.6} \begin{bmatrix} 11 \\ 14 - 30.88 \end{bmatrix} = \begin{bmatrix} 0.5934 \\ -0.8303 \end{bmatrix}$

For e_2 , $\begin{bmatrix} 0 \\ 0 \end{bmatrix} = (S - \lambda_2 I)u$
 $\begin{bmatrix} 14/3 - 6.15 & -11 \\ -11 & 23 - 6.15 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$
 $-11u_1 + (23 - 6.15)u_2 = 0$
 $-11u_1 + 16.85u_2 = 0$
 $\frac{u_1}{11} = \frac{u_2}{16.85}$
 $e_2 = \begin{bmatrix} 0.5934 \\ 0.8303 \end{bmatrix}$

5) $\begin{bmatrix} x_{1k} \\ x_{2k} \end{bmatrix}$
 $e_1^T \begin{bmatrix} x_{1k} - \bar{x}_1 \\ x_{2k} - \bar{x}_2 \end{bmatrix} = 0.5934(x_{1k} - \bar{x}_1) - 0.8303(x_{2k} - \bar{x}_2)$
 $= 0.5934(x_{1k} - \bar{x}_1) - 0.8303(x_{2k} - \bar{x}_2)$
 $= 0.5934(4-8) - 0.8303(11-8.5)$
 $= -4.305$

x_1	4	8	13	7
x_2	11	4	5	14

1st PCA -4.3052 | 5.2321 | 5.4128 | 5.1238

What is PCA
 It is a dimension reduction and ML method used to simplify the large data set into smaller set while maintaining significant patterns and trends.

Code:

```
from google.colab import files
heart=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy import stats
import seaborn as sns
from sklearn.preprocessing import LabelEncoder,
OneHotEncoder from sklearn.model_selection import
train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA

df1=pd.read_csv("heart (1).csv")
df1.head()
text_cols = df1.select_dtypes(include=['object']).columns
label_encoder = LabelEncoder()
for col in text_cols:
    df1[col] =
label_encoder.fit_transform(df1[col])
print(df1.head())
X = df1.drop('HeartDisease', axis=1)
y = df1['HeartDisease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42) scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train) X_test =
scaler.transform(X_test)
# Support Vector Machine
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)
svm_predictions = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_predictions)
print(f"SVM Accuracy: {svm_accuracy}")

# Logistic Regression
lr_model = LogisticRegression(random_state=42)
lr_model.fit(X_train, y_train) lr_predictions =
lr_model.predict(X_test) lr_accuracy =
accuracy_score(y_test, lr_predictions)
print(f"Logistic Regression Accuracy: {lr_accuracy}")
```

```

# Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print(f"Random Forest Accuracy: {rf_accuracy}")

models = {
    "SVM": svm_accuracy,
    "Logistic Regression":
        lr_accuracy, "Random Forest":
        rf_accuracy
}

best_model = max(models, key=models.get)
print(f"\nBest Model: {best_model} with accuracy {models[best_model]}")
pca = PCA(n_components=0.95)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

svm_model_pca = SVC(kernel='linear', random_state=42)
svm_model_pca.fit(X_train_pca, y_train)
svm_predictions_pca = svm_model_pca.predict(X_test_pca)
svm_accuracy_pca = accuracy_score(y_test, svm_predictions_pca)
print(f"SVM Accuracy (with PCA): {svm_accuracy_pca}")

lr_model_pca = LogisticRegression(random_state=42)
lr_model_pca.fit(X_train_pca, y_train)
lr_predictions_pca = lr_model_pca.predict(X_test_pca)
lr_accuracy_pca = accuracy_score(y_test, lr_predictions_pca)
print(f"Logistic Regression Accuracy (with PCA): {lr_accuracy_pca}")

rf_model_pca = RandomForestClassifier(random_state=42)
rf_model_pca.fit(X_train_pca, y_train)
rf_predictions_pca = rf_model_pca.predict(X_test_pca)
rf_accuracy_pca = accuracy_score(y_test, rf_predictions_pca)
print(f"Random Forest Accuracy (with PCA): {rf_accuracy_pca}")

models_pca = {
    "SVM": svm_accuracy_pca,
    "Logistic Regression": lr_accuracy_pca,
    "Random Forest": rf_accuracy_pca
}

best_model_pca = max(models_pca, key=models_pca.get)
print(f"\nBest Model (with PCA): {best_model_pca} with accuracy {models_pca[best_model_pca]}")

```